







(a)-(c) H200 SXM (NVSwitch), single node. Transpose and stencil are Parallel Research Kernels benchmarks; (c) whiskers span min to max over the 5 runs.

(b) At 16384^2 , WinIPC buffered reaches 8.8x host-staged MPI, within 1.4% of GPU-aware MPI.

(d) GB200 NVL, 4 and 8 nodes, over a cross-node CUDA fabric-handle window. WinIPC bars are the buffered variant; its single-kernel direct variant reaches 5.3x at 16 GPUs and 3.2x at 32 GPUs on transpose.