







(a)-(c) H200 SXM (NVSwitch), single node. (c) whiskers span min to max over the 5 runs.

(b) $B += A^T$, the Parallel Research Kernels operation the paper reports. At 16384^2 : WinIPC buffered is 8.8x host-staged MPI and 1.4% behind GPU-aware MPI.

(d) GB200 NVL scale-out system, 4 and 8 nodes, cross-node CUDA fabric-handle window. WinIPC = buffered variant, as in (b) — its single-kernel direct variant reaches 5.3x (16 GPUs) and 3.2x (32 GPUs) on transpose.