







(a)-(c) H200 SXM (NVSwitch), single node. (c) bars are medians of 5 runs, whiskers min/max; (a), (b) and (d) are single runs per point.

(d) GB200 NVL scale-out system, 4 and 8 nodes, cross-node CUDA fabric-handle window. WinIPC = buffered variant, as in (b) — its direct variant reaches 5.3x (16 GPUs) and 3.2x (32 GPUs) on transpose.

Stencil GPU-aware MPI is omitted from (d): those runs used zero warmup, so one-time lazy connection setup was charged to that variant alone. (d) is two configurations, not a scaling curve.