



(a) Accumulate, $B += A^T$ (b) Overwrite, $B = A^T$ 

Job 90161, 4 GPUs on one H200 node, 13 matrix orders, medians of 3 reps, 100 timed iterations after 20 untimed.
 Non-power-of-two orders are included so the single-kernel turnover is resolved rather than interpolated between powers of two.
 NVSHMEM was measured but is not plotted: its rates in this job disagree with job 63328 by up to 45x and are unexplained.
 Rates follow the paper's Eq. 1 ($2N^2 \cdot \text{sizeof}(\text{double})/t$), which counts the same bytes for both operations although accumulate must read the destination first, so (a) and (b) are not comparable as bandwidths.